

BTF

Contents

BTF

1. BTF:

1.1. 🧠 What is BTF?

1.1.1. 🧬 Why is BTF Important?

1.1.2. ⚙️ Use Cases of BTF

1.1.3. 🔧 How BTF Works

1.1.4. ✅ Summary

1.2. BTF Support (recompiling the kernel)?

1.2.1. What kernel support is needed for BTF?

1.2.2. What if your kernel doesn't have BTF?

1.2.3. Summary

1.2.4. TL;DR

1.3. bpftool:

1.3.1. 🔧 What is bpftool?

1.3.2. 🔍 What is bpftool Used For?

1.3.3. 🧐 How to Use bpftool to Read BTF Info

1.3.4. 💡 Practical Example

1.3.5. 📦 Installing bpftool

1.3.6. ✅ Summary

1.4. Example how bpftool and bpftrace

1.4.1. 🧴 1. CO-RE BPF Program Example (Trace Process Execs)

1.4.2. 🔧 Requirements

1.4.3. 📁 File: trace_exec.c

1.4.4. 🏗️ Compile with BTF

1.4.5. 🔍 2. How BTF Shows Up in bpftool

1.4.6. 🔍 View kernel BTF types:

1.4.7. 🔥 3. Run It with libbpf-based loader or bpftool

1.5. 🧠 4. BTF in bpftrace

1.5.0.1. Example:

1.5.1. ✅ Summary

1.6. Working Makefile or how to set up a vmlinux.h generation script?

1.6.1. 📄 Notes

2. Aya Rust based eBPF Development:

2.1. 🧠 What is Aya?

2.2. 📌 How the Above Concepts Fit with Aya

2.3. 🏠 CO-RE in Aya

2.3.1. 🔧 How Aya Builds and Loads eBPF

2.3.1.1. BPF Program (Rust) → CO-RE ELF

2.3.1.2. Loader → Loads BPF into kernel

2.3.2. 🔍 Checking BTF in Aya-generated ELF

2.3.3. ✅ Summary: How it All Fits Together with Aya

2.4. Example **Rust/Aya version of trace_exec**

3. Solana:

1. BTF:

1.1. 🧠 What is BTF?

BTF = **BPF Type Format**.

It's a **metadata format** used by the *eBPF subsystem* in the Linux kernel to describe **types, structures, and debugging information** for BPF programs – very similar to how DWARF works for traditional compiled binaries.

Think of BTF as a **lightweight debugging/type info system** for BPF programs.

1.1.1. 🧩 Why is BTF Important?

Modern **eBPF** programs interact with kernel structures and data, which:

- can vary across kernel versions
- are not always visible at runtime

BTF enables these programs to:

1. Understand and access kernel data structures at runtime.
2. Be **portable** across kernel versions (via CO-RE).
3. Allow tools like `bpftool` or `bpfttrace` to show symbols, fields, types.
4. Reduce the need for custom kernel headers.

1.1.2. ⚙️ Use Cases of BTF

<i>Use Case</i>	<i>Description</i>
CO-RE (Compile Once - Run Everywhere)	Enables portable eBPF programs that adapt to different kernel versions without recompilation.
bpftool introspection	Lets tools like bpftool show type info, structure layouts, and BPF maps/programs clearly.
bpfftrace/DTrace-like tools	Allows high-level tracing tools to generate field-safe probes dynamically.
Verifier Assistance	Helps the BPF verifier with type safety, pointer validation, etc.

1.1.3. 🔧 How BTF Works

- The kernel may expose a file:

```
/sys/kernel/btf/vmlinux
```

Which contains BTF metadata for the running kernel.

- When you compile a BPF program with clang, you can generate BTF data with:
`clang -target bpf -g -O2 -g -Xclang -target-feature -Xclang +btf ...`
- Tools like bpftool, libbpf, and llvm use this BTF section to:
 - read symbol types
 - access struct fields
 - perform CO-RE relocations

1.1.4. ✅ Summary

BTF = Metadata for BPF Programs

It allows:

- eBPF programs to understand kernel structure layouts
- runtime portability (CO-RE)
- user-friendly introspection and debugging
- better type safety in the kernel verifier

Without BTF, writing portable or introspectable eBPF programs would require building per-kernel versions or parsing headers manually – error-prone and fragile.

1.2. BTF Support (recompiling the kernel)?

Usually, yes – but it depends on your kernel version and distribution.

1.2.1. What kernel support is needed for BTF?

1. Kernel Config Options

To have full BTF support, the kernel must be built with these options enabled:

- `CONFIG_DEBUG_INFO_BTF=y` (This enables BTF debug info generation during kernel build.)
- `CONFIG_BPF_SYSCALL=y` (Enables BPF syscall support, usually already enabled.)
- `CONFIG_DEBUG_INFO=y` (Kernel debug info support; often needed together with BTF.)
- Optionally: `CONFIG_BPF_JIT_ALWAYS_ON` (for better BPF performance)

2. Kernel Version

- **Linux kernel 4.18+** started BTF support.
- **Kernel 5.3+** improved BTF and CO-RE support significantly.
- Many distributions backport BTF to their kernels.

3. BTF Data Availability

- After building the kernel with `CONFIG_DEBUG_INFO_BTF=y`, the **BTF data is embedded** inside the kernel img.
- This BTF data is typically exported at runtime as:

`/sys/kernel/btf/vmlinux`

- Tools (like `bpftool`) read this file to get kernel type info.

1.2.2. What if your kernel doesn't have BTF?

- You **won't** have `/sys/kernel/btf/vmlinux`.
- You can **generate and supply your own BTF** files but it's complex.
- Some tools might still work but with reduced functionality.
- Many modern distros now ship kernels with BTF enabled by default.

1.2.3. Summary

<i>Scenario</i>	<i>BTF availability</i>
Kernel built with <code>CONFIG_DEBUG_INFO_BTF=y</code>	Full BTF support + <code>/sys/kernel/btf/vmlinux</code> present
Kernel without BTF config	No BTF, limited eBPF capabilities
Older kernels (before 4.18)	No BTF support

1.2.4. TL;DR

For proper BTF support, **yes, you typically need to recompile the kernel with `CONFIG_DEBUG_INFO_BTF=y`**. If you use a modern distro kernel (5.x+), it's often enabled already.

1.3. bpftool:

It's a modern eBPF tooling, its a tool for inspection and simple manipulation of eBPF programs and maps.

1.3.1. 🛠️ What is bpftool?

bpftool an official utility provided by the Linux kernel for *managing, introspecting, and debugging BPF objects*, including:

- BPF programs
- Maps
- Links
- BTF (BPF Type Format) metadata
- cgroup attachments
- netlink and tc hooks

It comes with the **kernel source** (under tools/bpf/bpftool) and also packaged by most major Linux distros.

1.3.2. 🔍 What is bpftool Used For?

Function	Description
bpftool prog show	Lists all loaded BPF programs
bpftool map show	Shows loaded BPF maps
bpftool btf dump	Dumps BTF (BPF Type Format) metadata
bpftool feature probe	Probes kernel for supported eBPF features
bpftool net	Inspects network stack hooks (tc, xdp)
bpftool cgroup	Shows BPF programs attached to cgroups

1.3.3. 📖 How to Use bpftool to Read BTF Info

1. 🧠 Check if kernel BTF info is available

```
ls /sys/kernel/btf/vmlinux
```

If the file exists, the kernel exposes its BTF data (✅ good!).

2. 📖 Dump BTF types from the kernel

```
sudo bpftool btf dump file /sys/kernel/btf/vmlinux format c
```

This prints out the kernel's types in C-like syntax, e.g.:

```
struct task_struct {  
    ...  
    struct mm_struct *mm;  
    int pid;  
    ...  
};
```

You can also filter specific types:

```
sudo bpftool btf dump file /sys/kernel/btf/vmlinux format c | grep -A 20 "struct task_struct"
```

3. 🖋️ Inspect BTF in an ELF/BPF Object File

If you compiled a BPF program with `clang -g -target bpf`, the BTF info is embedded in the object file:

```
bpftool btf dump file myprog.o format c
```

This is useful for:

- Debugging structure layouts
- Verifying CO-RE compatibility
- Understanding what the verifier sees

1.3.4. 💡 Practical Example

```
bpftool btf dump file /sys/kernel/btf/vmlinux format raw
```

This dumps the BTF types in raw format (for scripting or analysis).

Or for human-readable output:

```
bpftool btf dump file /sys/kernel/btf/vmlinux format c
```

—

1.3.5. 🛒 Installing bpftool

- Ubuntu/Debian:

```
sudo apt install bpftool
```

- On Fedora:

```
sudo dnf install bpftool
```

Build from kernel source:

```
cd /usr/src/linux/tools/bpf/bpftool
```

```
make
```

```
sudo make install
```

1.3.6. ✅ Summary

Command	Purpose
<code>bpftool btf dump file /sys/kernel/btf/vmlinux format c</code>	View kernel BTF types
<code>bpftool btf dump file prog.o format c</code>	Inspect BTF in a BPF object
<code>bpftool prog/map/show</code>	List loaded BPF programs or maps
<code>bpftool feature probe</code>	See kernel eBPF/BTF capabilities

1.4. Example how bpftool and bpftrace

1.4.1. 1. CO-RE BPF Program Example (Trace Process Execs) real example of a CO-RE (Compile Once - Run Everywhere)

BPF program and how **BTF** shows up in bpftool and bpftrace.

This example traces `execve` calls using eBPF and accesses kernel struct fields via BTF.

1.4.2. Requirements

- Kernel 5.3+ with BTF enabled (`/sys/kernel/btf/vmlinux` exists)
- clang, llvm, bpftool, libbpf-dev
- BPF CO-RE headers (from kernel or linux-headers)

1.4.3. File: `trace_exec.c`

```
// SPDX-License-Identifier: GPL-2.0
#include "vmlinux.h"
#include <bpf/bpf_helpers.h>
#include <bpf/bpf_core_read.h>

char LICENSE[] SEC("license") = "GPL";

SEC("tracepoint/syscalls/sys_enter_execve")
int trace_exec(struct trace_event_raw_sys_enter* ctx)
{
    struct task_struct *task = (struct task_struct *)bpf_get_current_task();
    pid_t pid = BPF_CORE_READ(task, pid);
    char comm[16];
    BPF_CORE_READ_STR(&comm, task, comm);

    bpf_printk("exec: %d %s\n", pid, comm);
    return 0;
}
```

This uses:

- `BPF_CORE_READ()` - to safely access `task_struct::pid` across kernels.
- **No kernel headers** are needed - it reads types from **BTF** at runtime.

1.4.4. Compile with BTF

```
clang -g -O2 -target bpf \
-D__TARGET_ARCH_x86 \
-I/path/to/vmlinux.h \
-I. \
-c trace_exec.c -o trace_exec.o
```

You must have `vmlinux.h` (BTF-generated header). You can generate it using `bpftool`:

```
bpftool btf dump file /sys/kernel/btf/vmlinux format c > vmlinux.h
```

Or:

```
bpftool btf dump file /sys/kernel/btf/vmlinux format c | pahole -J - > vmlinux.h
```


1.4.5. 🔍 2. How BTF Shows Up in bpftool

✳️ Inspect BTF info in a BPF object:

```
bpftool btf dump file trace_exec.o format c
```

You'll see:

```
struct task_struct {
    int pid;
    char comm[16];
    ...
};
```

This proves the object has **embedded BTF metadata**.

1.4.6. 🔍 View kernel BTF types:

```
bpftool btf dump file /sys/kernel/btf/vmlinux format c | grep -A 5 'task_struct'
```

You'll see the actual kernel task_struct layout, used by CO-RE.

—

1.4.7. 🚒 3. Run It with libbpf-based loader or bpftool

If using bpftool to load:

```
sudo bpftool prog load trace_exec.o /sys/fs/bpf/trace_exec
```

Attach it:

```
sudo bpftool prog attach /sys/fs/bpf/trace_exec \
    tracepoint syscalls sys_enter_execve
```

Then check the logs:

```
sudo cat /sys/kernel/debug/tracing/trace_pipe
```

You'll see:

```
exec: 12345 bash
exec: 12346 ls
```

1.5. 🧠 4. BTF in bpftrace

1.5.0.1. Example:

```
sudo bpftrace -e 'kprobe:do_exit { printf("Exiting PID: %d\n", pid); }'
```

bpftrace automatically uses **BTF** from /sys/kernel/btf/vmlinux to know what do_exit is and how to access its parameters or fields.

It enables zero-effort tracing using BTF introspection.

1.5.1. ✅ Summary

Tool	How BTF Helps
bpftool	Dumps struct/type info from BPF objs or kernel
bpftrace	Auto-detects kernel struct layouts via BTF
libbpf	Enables BPF CO-RE with BTF relocations
clang	Emits .BTF section in object files

:

1.6. Working Makefile or how to set up a vmlinux.h generation script?

Everything you need to **compile and load a BPF CO-RE program**, including a **Makefile** and a way to **generate vmlinux.h** automatically using bpftool.

✓ Folder Structure

```
bpf-core-trace/
├─ Makefile
├─ trace_exec.c
└─ vmlinux.h      <-- auto-generated
```

1. trace_exec.c (same as before)

```
// SPDX-License-Identifier: GPL-2.0
#include "vmlinux.h"
#include <bpf/bpf_helpers.h>
#include <bpf/bpf_core_read.h>

char LICENSE[] SEC("license") = "GPL";

SEC("tracepoint/syscalls/sys_enter_execve")
int trace_exec(struct trace_event_raw_sys_enter* ctx)
{
    struct task_struct *task = (struct task_struct *)bpf_get_current_task();
    pid_t pid = BPF_CORE_READ(task, pid);
    char comm[16];
    BPF_CORE_READ_STR(&comm, task, comm);

    bpf_printk("exec: %d %s\n", pid, comm);
    return 0;
}
```

2. Makefile

```
CLANG ?= clang
BPFTOOL ?= bpftool
BPF_OBJ = trace_exec.o
SRC = trace_exec.c

.PHONY: all clean

all: $(BPF_OBJ)

# Generate vmlinux.h from kernel BTF
vmlinux.h:
    @if [ ! -e /sys/kernel/btf/vmlinux ]; then \
        echo "ERROR: Kernel BTF not found at /sys/kernel/btf/vmlinux"; \
        exit 1; \
    fi
    @echo "[*] Generating vmlinux.h from BTF..."
    $(BPFTOOL) btf dump file /sys/kernel/btf/vmlinux format c > vmlinux.h

$(BPF_OBJ): $(SRC) vmlinux.h
    @echo "[*] Compiling BPF program..."
    $(CLANG) -g -O2 -target bpf \
```

```
-D__TARGET_ARCH_$(shell uname -m | sed 's/x86_64/x86/') \
-I. \
-c $(SRC) -o $(BPF_OBJ)
```

clean:

```
rm -f $(BPF_OBJ) vmlinux.h
```

3. Build the BPF Program

make

This will:

- Check for /sys/kernel/btf/vmlinux
- Generate vmlinux.h
- Compile trace_exec.o with BTF info

4. Load & Attach the BPF Program

You can load and attach it using bpftool or a loader like [libbpf](https://github.com/libbpf/libbpf).

Using bpftool:

```
# Load the BPF program
```

```
sudo bpftool prog load trace_exec.o /sys/fs/bpf/trace_exec type tracepoint
```

```
# Attach it to tracepoint
```

```
sudo bpftool prog attach name trace_exec \
    tracepoint syscalls:sys_enter_execve
```

5. See the Output

Read from trace_pipe:

```
sudo cat /sys/kernel/debug/tracing/trace_pipe
```

You'll see something like:

```
exec: 2345 bash
```

```
exec: 2346 ls
```

1.6.1. Notes

- If you get vmlinux.h: file not found – ensure /sys/kernel/btf/vmlinux exists.
- Requires: kernel 5.4+ (for stable CO-RE), bpftool, clang, llvm, and BPF headers.

2. Aya Rust based eBPF Development:

Connecting the above with **Aya** a **Rust-based eBPF development**.

Aya is a modern, pure-Rust framework for writing and loading eBPF programs **without relying on libbpf or C code**.

2.1. 🧠 What is Aya?

Aya is:

- A **Rust framework** for writing eBPF programs (like `trace_exec`) in Rust
- Provides **CO-RE** (Compile Once - Run Everywhere) support using **BTF**
- Does **not require libbpf or clang/LLVM** (only Rust toolchain)
- Works on **modern Linux kernels** (with BTF)

> It bridges **Rust ↔ Kernel BPF**, using BTF for safe type access like the `BPF_CORE_READ()` macro in C, but with type-safe Rust code.

2.2. 📌 How the Above Concepts Fit with Aya

Concept	In C (libbpf)	In Rust (Aya)
Program language	C	Rust
Loader	bpftool / custom C	Rust user-space app with Aya runtime
Kernel struct access	<code>BPF_CORE_READ()</code> + <code>vmlinux.h</code>	Aya auto-generates BTF access
BTF requirement	<code>/sys/kernel/btf/vmlinux</code>	✅ Required for CO-RE (via BTF)
Program sections	<code>SEC("tracepoint/...")</code>	<code>#[map]</code> , <code>#[tracepoint]</code> , etc. macros
Compilation	<code>clang -target bpf</code>	<code>cargo xtask</code> (or <code>cargo bpf</code>)
Safety	Unsafe C	Safe or unsafe Rust

2.3. 🔄 CO-RE in Aya

Just like in the C example where we used:

```
BPF_CORE_READ(task, pid);
```

Aya lets you **generate and reference BTF types from Rust** using macros like:

```
use aya_bpf::bindings::task_struct;

#[tracepoint(name = "sys_enter_execve")]
pub fn trace_exec(ctx: TracePointContext) -> i32 {
    unsafe {
        let task: *const task_struct = bpf_get_current_task();
        let pid = (*task).pid;
        // ... log or trace comm
    }
    0
}
```

Aya uses **bindgen** and **BTF introspection** behind the scenes to safely access those fields.

2.3.1. 🛠️ How Aya Builds and Loads eBPF

Aya typically uses a split-project layout:

```
project/
├─ xtask/           # Loader/CLI
├─ bpf/             # eBPF programs in Rust
│   └─ src/
│       └─ trace_exec.rs
└─ Cargo.toml
```

2.3.1.1. BPF Program (Rust) → CO-RE ELF

The eBPF program is compiled using:

```
cargo xtask build-ebpf
```

It embeds:

- .BTF section for CO-RE
- Symbol section (like in `trace_exec.o`)
- Maps and program sections

2.3.1.2. Loader → Loads BPF into kernel

```
let mut bpf = Bpf::load_file("target/bpf/trace_exec.o"?);
program.attach()?;
```

The loader handles:

- BTF relocations
- Program/map loading
- Attachments to hooks

2.3.2. 🔍 Checking BTF in Aya-generated ELF

You can inspect BTF in an Aya-generated `.o` with `bpftool`, just like C:

```
bpftool btf dump file target/bpf/trace_exec.o format c
```

You'll see the Rust-generated types used in the BPF program – verified via BTF.

2.3.3. ✅ Summary: How it All Fits Together with Aya

Component	Role
BTF	Provides kernel type metadata
CO-RE	Lets programs access structs like <code>task_struct</code> across kernel versions
Aya	Rust framework for writing + loading BPF
<code>trace_exec.c</code>	Equivalent written in Aya's Rust DSL
<code>vmlinux.h</code>	Not used directly – Aya uses BTF instead
<code>bpftool</code>	Still useful for inspecting objects

2.4. Example Rust/Aya version of trace_exec

Minimal Aya Rust CO-RE eBPF example that traces `execve` syscalls just like the C example, with instructions and a sample project structure.

Aya CO-RE Example: Trace `execve`

Step 1. Project structure

```
aya-trace-exec/  
├── bpf/  
│   ├── Cargo.toml  
│   └── src/  
│       └── trace_exec.rs  
├── Cargo.toml  
└── src/  
    └── main.rs          # Loader program
```

Step 2. `bpf/Cargo.toml` – eBPF program crate

```
[package]  
name = "trace_exec"  
version = "0.1.0"  
edition = "2021"  
  
[dependencies]  
aya-bpf = { version = "0.19", features = ["tracepoint"] }  
aya-log = "0.19"  
  
[lib]  
crate-type = ["cdylib"]  
—
```

Step 3. `bpf/src/trace_exec.rs` – BPF program

```
#![no_std]  
#![no_main]  
  
use aya_bpf::{  
    bindings::task_struct,  
    macros::tracepoint,  
    programs::TracePointContext,  
    helpers::bpf_get_current_task,  
};  
use aya_log_ebpf::info;  
  
#[tracepoint(name = "sys_enter_execve")]  
pub fn trace_exec(ctx: TracePointContext) -> u32 {  
    match unsafe { try_trace_exec(ctx) } {  
        Ok(ret) => ret,  
        Err(_) => 1,  
    }  
}  
  
unsafe fn try_trace_exec(_ctx: TracePointContext) -> Result<u32, ()> {  
    let task = bpf_get_current_task();  
    if task.is_null() {  
        return Err(());  
    }  
}
```

```

let pid = (*task).pid;
// comm is fixed-size 16 bytes char array
let comm = (*task).comm;

// Log the exec call with PID and comm (command name)
// This shows up in `trace_pipe` if tracing is enabled
info!(&format!("exec: {} {:?}", pid, comm));

Ok(0)
}

```

Step 4. Top-level Cargo.toml

```

[workspace]
members = [
    "bpf",
]

```

Step 5. src/main.rs – Loader

```

use aya::{Bpf, programs::TracePoint, util::online_cpus};
use aya::maps::perf::PerfEventArray;
use aya::util::online_cpus;
use std::convert::TryInto;
use std::error::Error;
use std::sync::Arc;
use tokio::{signal, task};

#[tokio::main]
async fn main() -> Result<(), Box<dyn Error>> {
    // Load the compiled eBPF program
    let mut bpf = Bpf::load_file("bpf/target/bpfel-unknown-none/release/
trace_exec.o")?;

    // Attach tracepoint
    let program: &mut TracePoint = bpf.program_mut("trace_exec").unwrap().try_into()?;
    program.load()?;
    program.attach("syscalls", "sys_enter_execve")?;

    println!("Attached to sys_enter_execve tracepoint. Waiting for events...");

    // Keep the program running
    signal::ctrl_c().await?;
    println!("Exiting...");

    Ok(())
}

```

Step 6. Build instructions

You need Rust nightly with cargo and rustup set up.

- Install Rust and aya-bpf dependencies

```

rustup target add bpfel-unknown-none
cargo install cargo-binutils llvm-tools-preview

```

- Build the BPF program

```
cd bpf
cargo build --release --target bpfel-unknown-none
```

- Build and run the loader

```
cd ..
cargo run --release
```

Step 7. View Output

Open a separate terminal:

```
sudo cat /sys/kernel/debug/tracing/trace_pipe
```

You should see lines like:

```
exec: 12345 [bash, ...]
```

Notes:

- Aya uses BTF automatically – no need for `vmlinux.h`
- The `trace_exec.rs` uses safe Rust with some `unsafe` block to dereference kernel pointers
- This example uses `aya-log` for simple `bpf_printk` logging
- Aya supports other program types too – kprobes, uprobes, xdp, etc.

3. Solana:

- Solana smart contracts are mostly written in Rust, often compiled into BPF Bytecode.
- Experience writing CO-RE BPF programs in Rust (eg: Aya) is highly transferable to Solana contract devl.
- Note As Solana BPF is a custom, sandboxed VM for executing blockchain programs and Solana BPF VM does not use BTF and similar kernel debugging infrastructure.
- Solana programs have deterministic execution, limited compute budgets and no kernel interaction.